EXECUTIVE ASSESSMENT

AI Benchmark Report: Business Analytics & Dashboard Development

Comparative evaluation of LLM performance on Kenya multi-region retail sales analysis (Q1 2026).

Dataset Scope: 24 Transactions (Nairobi, Mombasa, Kisumu, Nakuru) **Date:** 9 November 2026 **Benchmark Standard:** Enterprise BI & Vanilla JS

Across all three models evaluated, the underlying analytical engine demonstrated complete mathematical precision by correctly recognizing KES 1,558,000 in revenue, KES 494,000 in gross profit (32.54% completed markup margin; 31.71% recognized gross margin), and 50 units sold from 21 completed orders, while properly excluding the KES 90,000 in phone returns and KES 50,000 in tablet cancellations. Because the core computational baseline was identical across the board, the benchmark differentiation rested heavily on user experience, information architecture, visual polish, and frontend execution. Composer 2.5 achieved a baseline score of 83.5/100, providing a functional but standard layout with basic KPI cards and Chart.js visualizations that lacked deeper UX refinement.

Grok 4.6 and Grok 4.7 elevated the assessment significantly, scoring 89.0/100 and 96.5/100, respectively. Grok 4.6 separated itself through a distinct executive-ledger design, utilizing a curated font palette (Sora, Source Sans 3, IBM Plex Mono), structured "Finding/Do" recommendation blocks, an isolated Exceptions audit table, and an intuitive "City Cash Position" bar meter. Grok 4.7 secured the top benchmark position by delivering enterprise-grade interactivity, replacing clunky dropdowns with touch-friendly quick-filter chips, introducing a sticky table-of-contents navigation bar, displaying an end-to-end transactional Order Ledger with audit totals, and maintaining robust dynamic recalculations across filtered subsets without losing global Q1 strategic context.

Benchmark Scoring & Attribute Matrix

MODEL / DELIVERABLE	MATH ACCURACY (35%)	AUDIT & DATA PREP (20%)	FRONTEND & UX (25%)	BI RIGOR (20%)	OVERALL SCORE
Composer 2.5 (Baseline BI)	35.0 / 35	15.0 / 20	18.0 / 25	15.5 / 20	**83.5 / 100**
Grok 4.6 (Ledger Style)	35.0 / 35	18.0 / 20	20.0 / 25	16.0 / 20	**89.0 / 100**
Grok 4.7 (Interactive Enterprise)	35.0 / 35	20.0 / 20	24.5 / 25	17.0 / 20	**96.5 / 100**

Architectural Highlights & Evaluation Takeaways

Grok 4.6: Executive Editorial Polish

Distinguished by strong visual hierarchy and a CFO-ready feel. The "City Cash Position" visual meter quickly anchors regional volume, while the "Finding vs. Do" structure provides clear, actionable takeaways. An isolated Exceptions audit table cleanly explains non-completed transactions.

Grok 4.7: Enterprise Reactive Architecture

Winner on technical frontend execution. Employs responsive chip-based filters, a persistent Table of Contents navigation bar, full order ledger reconciliation in the footer, and comprehensive ARIA live regions for accessibility and screen reader support.